

Advanced Tuning Methodologies

Dr.-Ing. Michael Klemm
Software and Services Group
Intel Corporation
(michael.klemm@intel.com)

Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Performance tests and ratings are measured using specific computer systems and/or components and reflect the approximate performance of Intel products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance. Buyers should consult other sources of information to evaluate the performance of systems or components they are considering purchasing. For more information on performance tests and on the performance of Intel products, reference www.intel.com/software/products.

Copyright © 2015, Intel Corporation. All rights reserved. Intel, the Intel logo, Xeon, Core, Xeon Phi, VTune, and Cilk are trademarks of Intel Corporation in the U.S. and other countries. *Other names and brands may be claimed as the property of others.

Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Shameless Plug...



Authors: Alexander Supalov, Andrey Semin, Michael Klemm, Chris Dahnken

Table of Contents:

Foreword by Bronis de Supinski (CTO LLNL)

Preface

Chapter 1: No Time to Read this Book?

Chapter 2: Overview of Platform Architectures

Chapter 3: Top-Down Software Optimization

Chapter 4: Addressing System Bottlenecks

Chapter 5: Addressing Application Bottlenecks:
Distributed Memory

Chapter 6: Addressing Application Bottlenecks:
Shared Memory

Chapter 7: Addressing Microarchitecture Bottlenecks

Chapter 8: Application Design Implications

ISBN-13 (pbk): 978-1-4302-6496-5

ISBN-13 (electronic): 978-1-4302-6497-2



Order now at <http://www.apress.com/9781430264965>



Software & Services Group, Developer Products Division

Copyright© 2015, Intel Corporation. All rights reserved. *Other brands and names are the property of their respective owners.

Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints





The Software Optimization Process

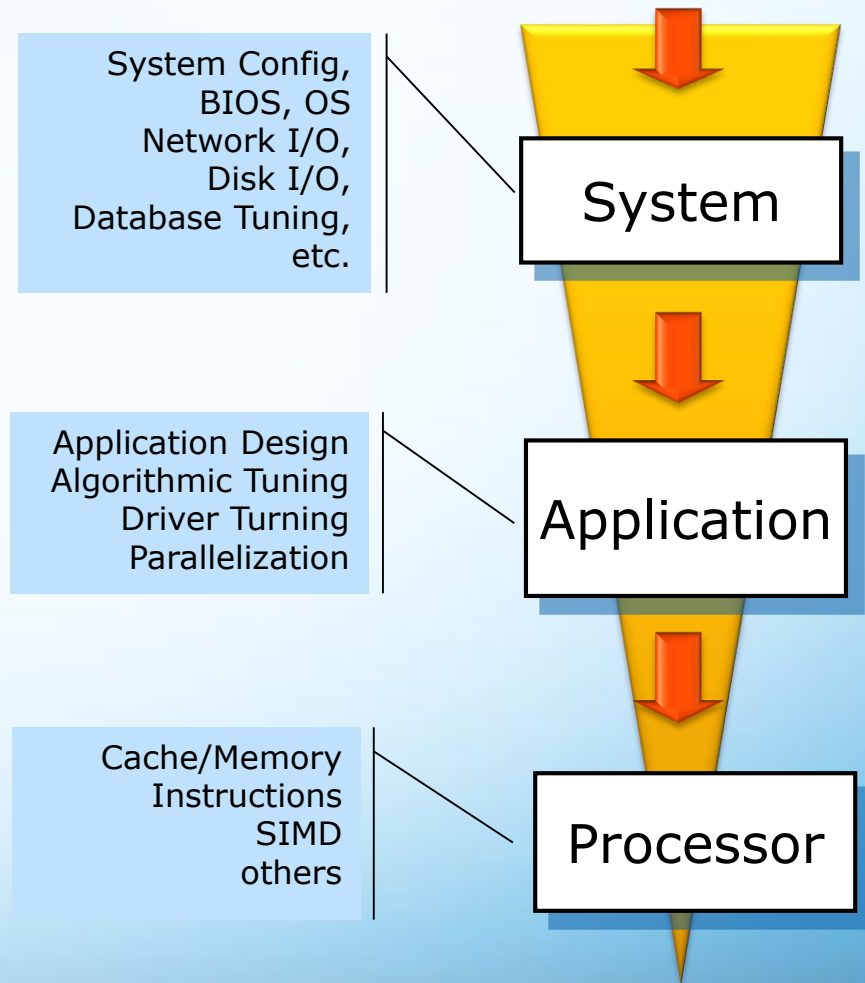
- The process of improving the software by eliminating bottlenecks so that it operates more efficiently on a given hardware and uses resources optimally
 - Identifying bottlenecks in the target application and eliminating them appropriately is the key to an efficient optimization
- There are many optimization methodologies, which help developers answer the questions of
 - Why to optimize?
 - What to optimize?
 - To what to optimize?
 - How to optimize?

These methods aid developers to reach their performance requirements.

Performance Analysis Methodology

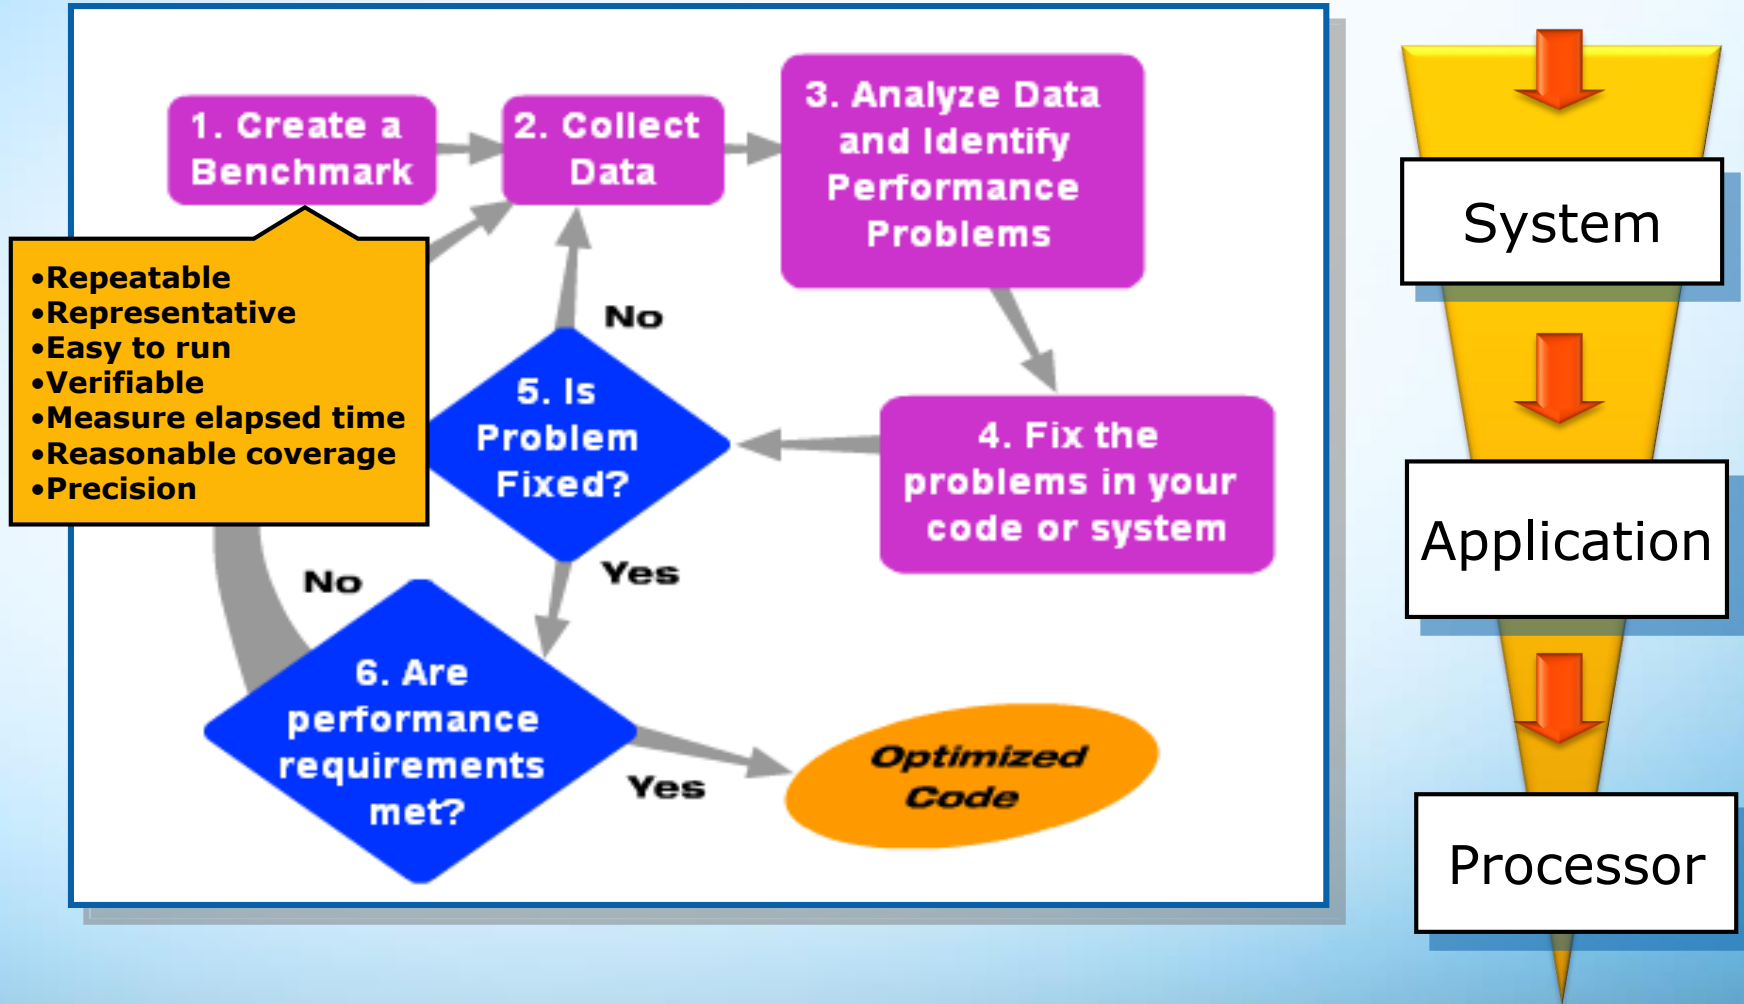
Optimization: A Top-down Approach

- Use top down approach
- Understand application and system characteristics
 - Use appropriate tools at each level



Performance Analysis Methodology

Optimization: A Top-down Approach



When to Stop

- Is architecture at maximum efficiency?
 - What this means: calculating theoretical maximum.
 - Know about best values for CPI or IPC.
 - Know the maximum FLOPS for the data type used.
- Is the performance requirement fulfilled?
 - What are the performance requirements?
 - Incrementally complete optimizations until done.
 - Key question: Are you “happy” with it?

CPI: Cycles per Instructions

IPC: Instructions per Cycle

FLOPS: Floating-Point Oper. Per Sec.

Questions to Ask Yourself

*"We should forget about small efficiencies, say about 97% of the time: **premature optimization is the root of all evil.**"*

— Donald Knuth


Quality code is:

- Portable
- Readable
- Maintainable
- Reliable


Intelligently sacrifice quality for performance

Amdahl's Law Slightly Reinterpreted


1 core, $t=140$



1 core, optimized



1 cores, optimized even more



A 2x improvement in the hotspots overall leads to 1.5x

A 10x improvement in the hotspots leads to 2.8x

Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Selecting type of data collection

tachyon_vtune_amp_xe - Microsoft Visual Studio

File Edit View Project Build Debug Tools Test Window Help

Release Win32

New Amplifier XE Result Start Page

Choose Analysis Type

Analysis Type

- Algorithm Analysis
 - Lightweight Hotspots
 - Hotspots
 - Concurrency
 - Locks and Waits
- Advanced Intel(R) Core(TM) 2 Processor
 - Bandwidth
 - Bandwidth Breakdown
 - Cycles and uOps
 - General Exploration
 - Memory Access
- Advanced Intel(R) Microarchitectured Core
 - Cycles and uOps
 - Front End Investigation
 - General Exploration
 - Memory Access
- Custom Analysis

Weight Hotspots

Identify your most time-consuming source code. Unlike Hotspots, Lightweight Hotspots does not collect stack information but can be used system-wide. This analysis type is based on the hardware event-based sampling collection. Press F1 for more details.

Details

Events configured for CPU: 45nm Intel(R) Core(TM) microarchi

NOTE: For analysis purposes, Amplifier XE may adjust the Sample After values in the table below by a multiplier. The multiplier depends on the value of the Duration time estimate.

After	Event
100000	Core cycles when
2000000	Reference cycles
2000000	Instructions retire

Start

Start Paused

Project Properties

Get Command Line

New ...

Copy from current

New CPU Event Collector analysis

New Stack-sampling Collector analysis

Callouts:

- All available analysis types
- Different ways to start the analysis
- Helps creating new analysis types
- Copy correct command line syntax to clipboard

VTune™ Amplifier XE

GUI Layout

Adjust Data Grouping

Function - Call Stack
Module - Function - Call Stack
Source File - Function - Call Stack
Thread - Function - Call Stack
... (Partial list shown)

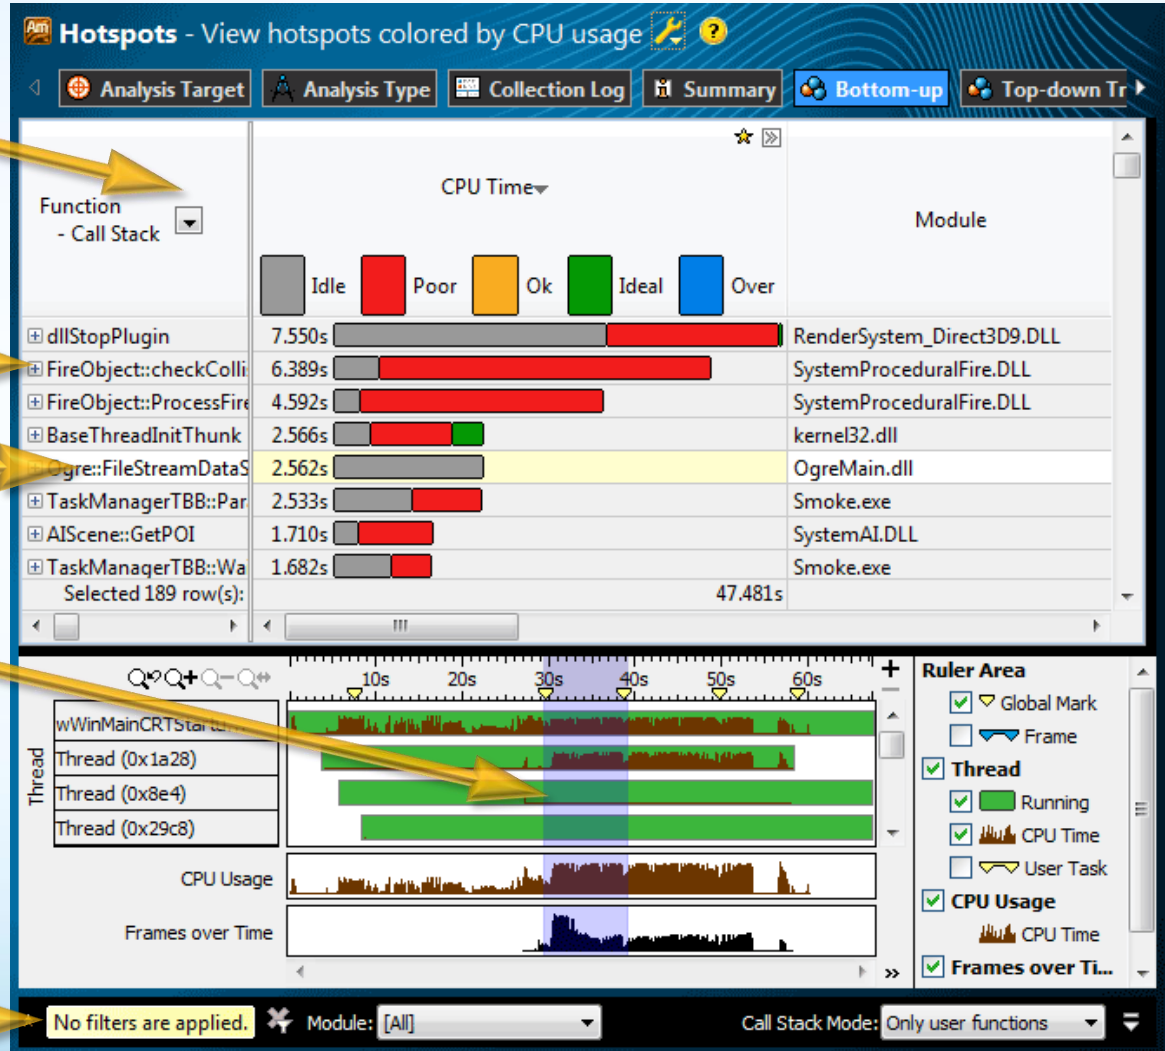
Click [+] for Call Stack

Double Click Function to View Source

Filter by Timeline Selection (or by Grid Selection)

Zoom In And Filter On Selection
Filter In by Selection
Remove All Filters

Filter by Module & Other Controls



VTune™ Amplifier XE

GUI Layout

Hotspots - View CPU time hotspots and stacks

Analysis Target Analysis Type Collection Log Summary Bottom-up Top-down Tree FireObj...

Line	Source	CPU Time	Address	Assembly	CPU Time
469	FireObject::checkCollision(V3 pos, V3 pre	0.476s	0x388c	fld st0, dword ptr [esp+0xc]	0.004s
472	#define FMax std::max<float>	0.561s	0x3890	fld st0, st0	0.993s
	- prev	6.846s	0x3892	fmlp st2, st0	0.787s
		3.593s	0x3894	fxch st0, st1	1.465s
477	float param2 = (AABB.zMax - prevPos.	0.830s	0x3896	fstp dword ptr [esp+0x8], st0	0.325s
478	bool neg = (rz < 0.f);	0.615s	0x389a	fld st0, dword ptr [esp+0x40]	0.014s
479	minP = FMax(neg? param2 : param1, mi	3.008s	0x389e	fsubrp st2, st0	0.010s
480	maxP = FMin(neg? param1 : param2, ma	1.875s	0x38a0	fld st0, st0	0.010s
481	if(maxP > minP) {	0.972s	0x38a2	fmlp st2, st0	0.233s
482	rx = 1.f/(pos.x - prevPos.x);	0.252s	0x38a4	fxch st0, st1	0.247s
483	param1 = (AABB.xMin - prevPos.x)	0.264s	0x38a6	fstp dword ptr [esp+0xc], st0	0.326s
	Pos.x)	0.040s	0x38aa	fcomp st0, st2	0.032s
	param1	0.047s	0x38ac	fnstsw ax	
	param2	0.274s	0x38ae	test ah, ax	
		0.164s	0x38b1	jp 0x100038c1	
		0.612s	0x38b3	Block 2:	
			0x38b3	mov dl, 0x1	
			0x38b5	lea ecx, ptr [esp+0xc]	0.024s
			0x38b9	jmp 0x100038c1 <Block 4>	
			0x38bb	Block 3:	
			0x38bb	xor dl, dl	0.159s

Selected 1 row(s): 0.830s Highlighted 6 row(s): 0.830s

Intel® VTune™ Amplifier XE

Time on Source / Asm

Quick Asm navigation:
Select source to highlight Asm

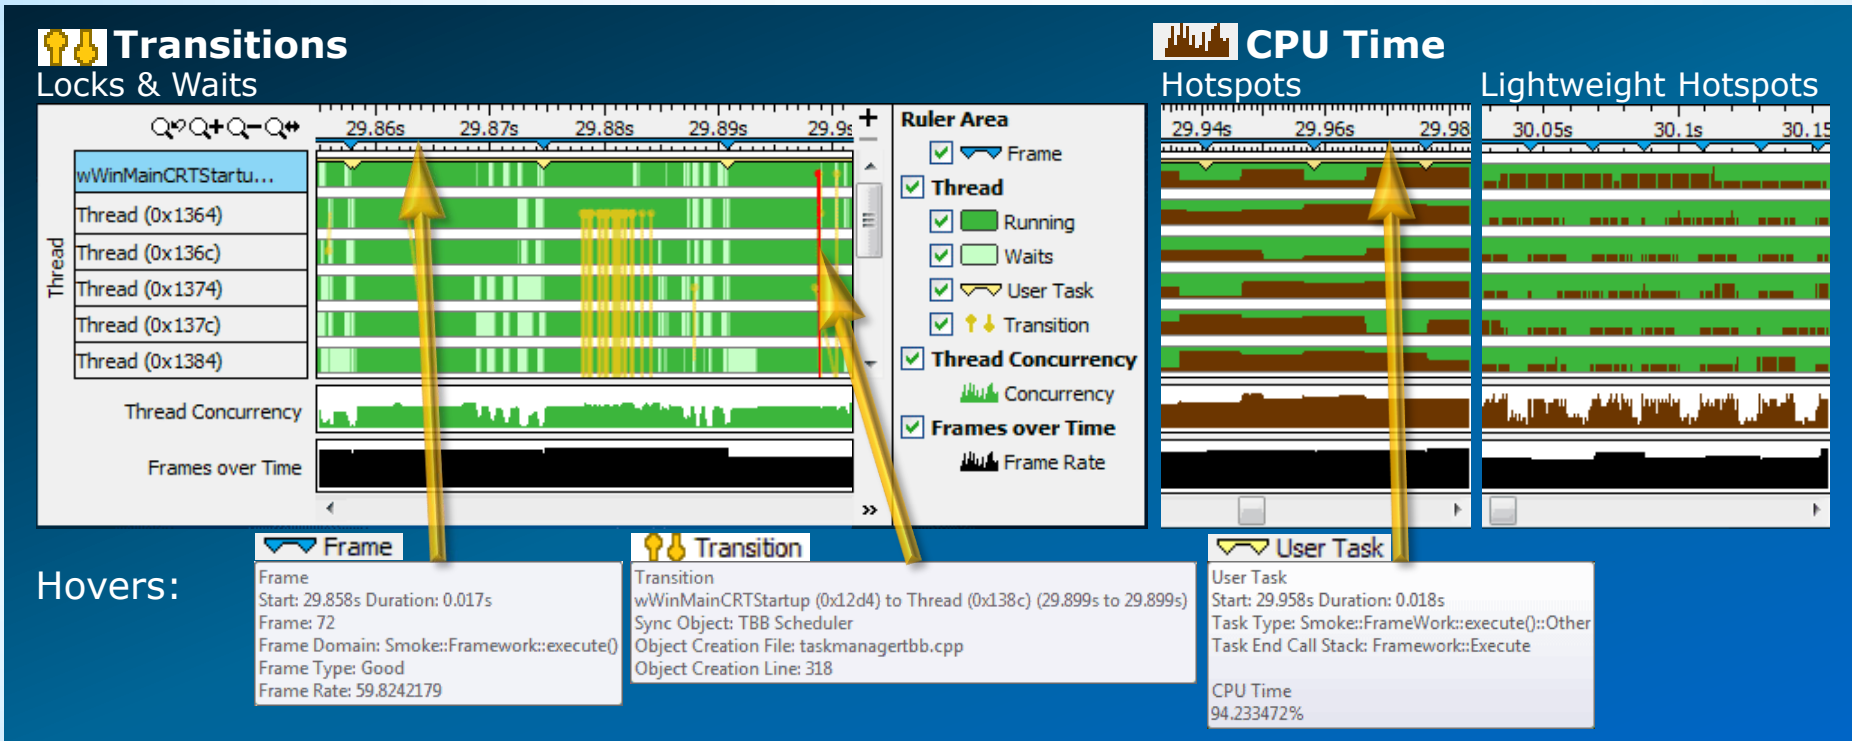
Quickly scroll to hot spots.
Scroll Bar "Heat Map" is an
overview of hot spots

Right click for instruction
reference manual

Click jump to scroll Asm

VTune™ Amplifier XE

GUI Layout



- Optional: Use API to mark frames and user tasks
- Optional: Add a mark during collection



Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Readying Your Application for Intel VTune Amplifier XE

- You should run Amplifier XE on a “Released/Optimized” build.
- Symbols Allow you to view the source (not just the assembly)
 - Windows: /Zi
 - Linux: -g
- Intel Threading Runtimes need instrumented runtimes
 - TBB: Define TBB_USE_THREADING_TOOLS
 - OpenMP: Use Intel Dynamic Version of OpenMP
- Call Stack Mode – Requires use of the dynamic version of the C Runtime library to properly attribute System Calls
 - Windows use: /MD(d)
 - Linux do not use: -static

Analysis Types

Basic Hotspots

- For each sample, capture execution context, time passed since previous sample and thread CPU time
- Allows time spent in system calls to be attributed to user functions making the call
- Provides additional knobs:
 - The defaults for Hotspot analysis are configurable and can be done so by creating a custom analysis type inherited from Hotspots

User-mode Sampling and Tracing Settings

CPU sampling interval, ms:	10	▲▼
Collect CPU sampling data:	With stacks	▼
Collect signalling API data:	No	▼
Collect synchronization API data:	No	▼
Collect I/O API data:	No	▼
<input checked="" type="checkbox"/> Collect timeline data		

Analysis Types

Advanced Hotspots

- Similar to Hotspot Analysis
 - Sampling is performed with the SEP collector
 - Driver is required
- Stack walking can be performed
 - Only hotspots are reported by default
 - More complex data collection is possible, incl. trip counts
- Samples are taken more frequently, but may have less accurate timing information
- Analysis may be performed for a single application or for the entire system

Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Issues in Parallel Applications

- Load imbalance
 - Work distribution is not optimal
 - Some threads are heavily loaded, while others idle
 - Slowest thread determines total speed-up
- Locking issues
 - Locks prohibit threads to concurrently enter code regions
 - Effectively serialize execution
- Parallelization overhead
 - With large no. of threads, data partition get smaller
 - Overhead might get significant (e.g. OpenMP startup time)


Issues in Parallel Applications

- Load imbalance
 - Work distribution is not optimal
 - Some threads are heavily loaded, while others idle
 - Slowest thread determines total speed-up
- Locking issues
 - Locks prohibit threads to concurrently enter code regions
 - Effectively serialize execution
- Parallelization overhead
 - With large no. of threads, data partition get smaller
 - Overhead might get significant (e.g. OpenMP startup time)

Threading Analysis Terminology

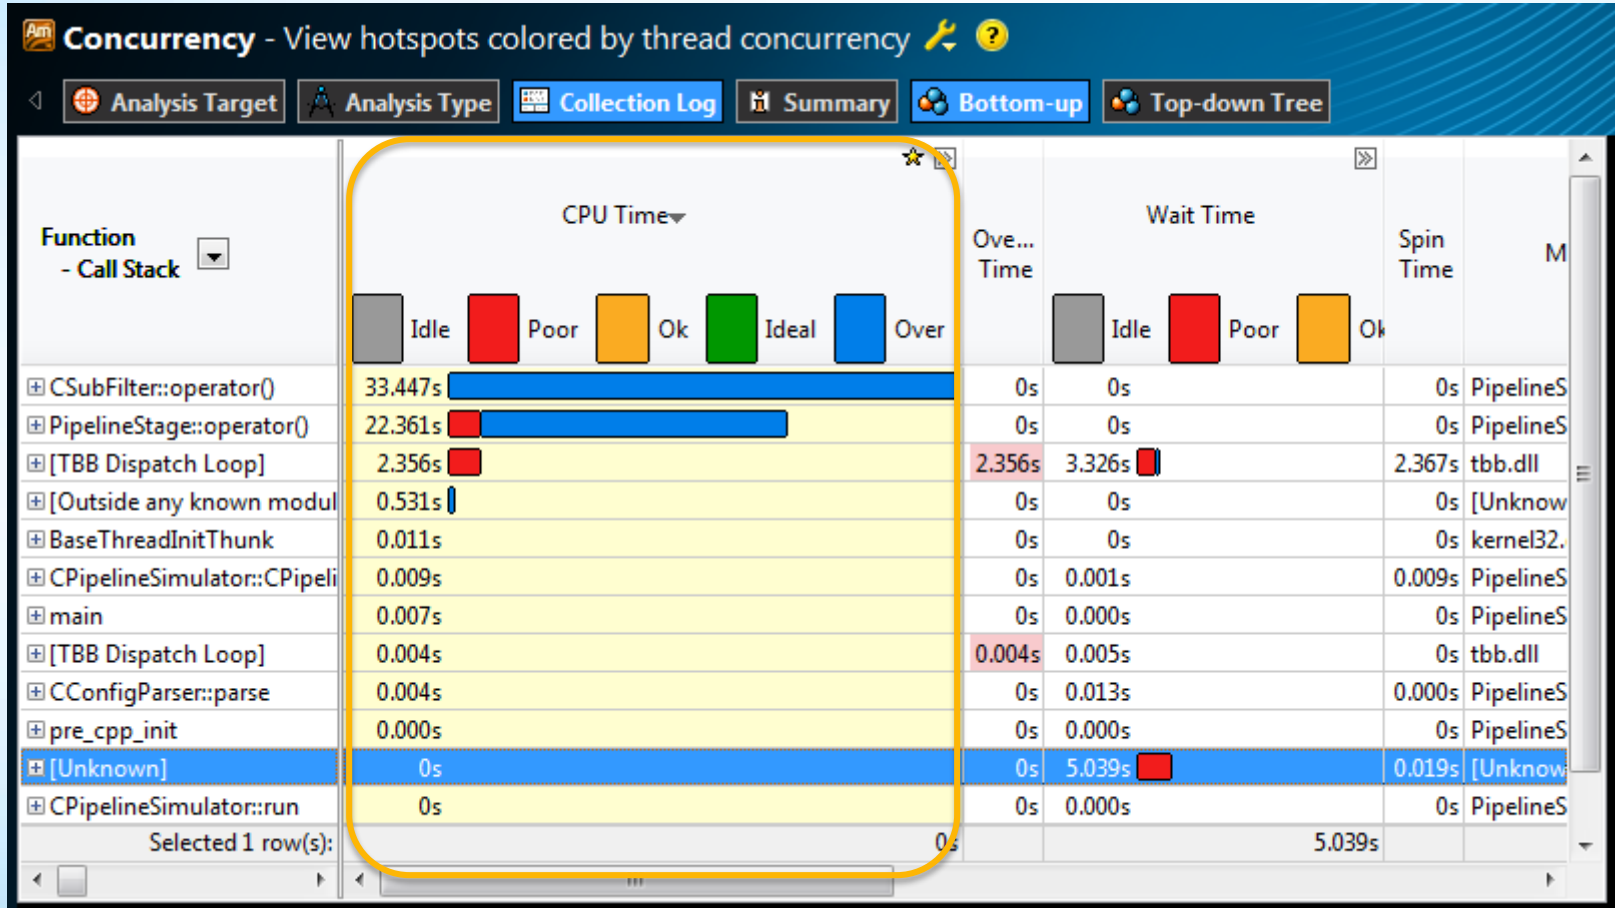


- **Elapsed Time:** 6 seconds
- **CPU Time:** T1 (4s) + T2 (3s) + T3 (3s) = 10 seconds
- **Wait Time:** T1(2s) + T2(2s) + T3 (2s) = 6 seconds



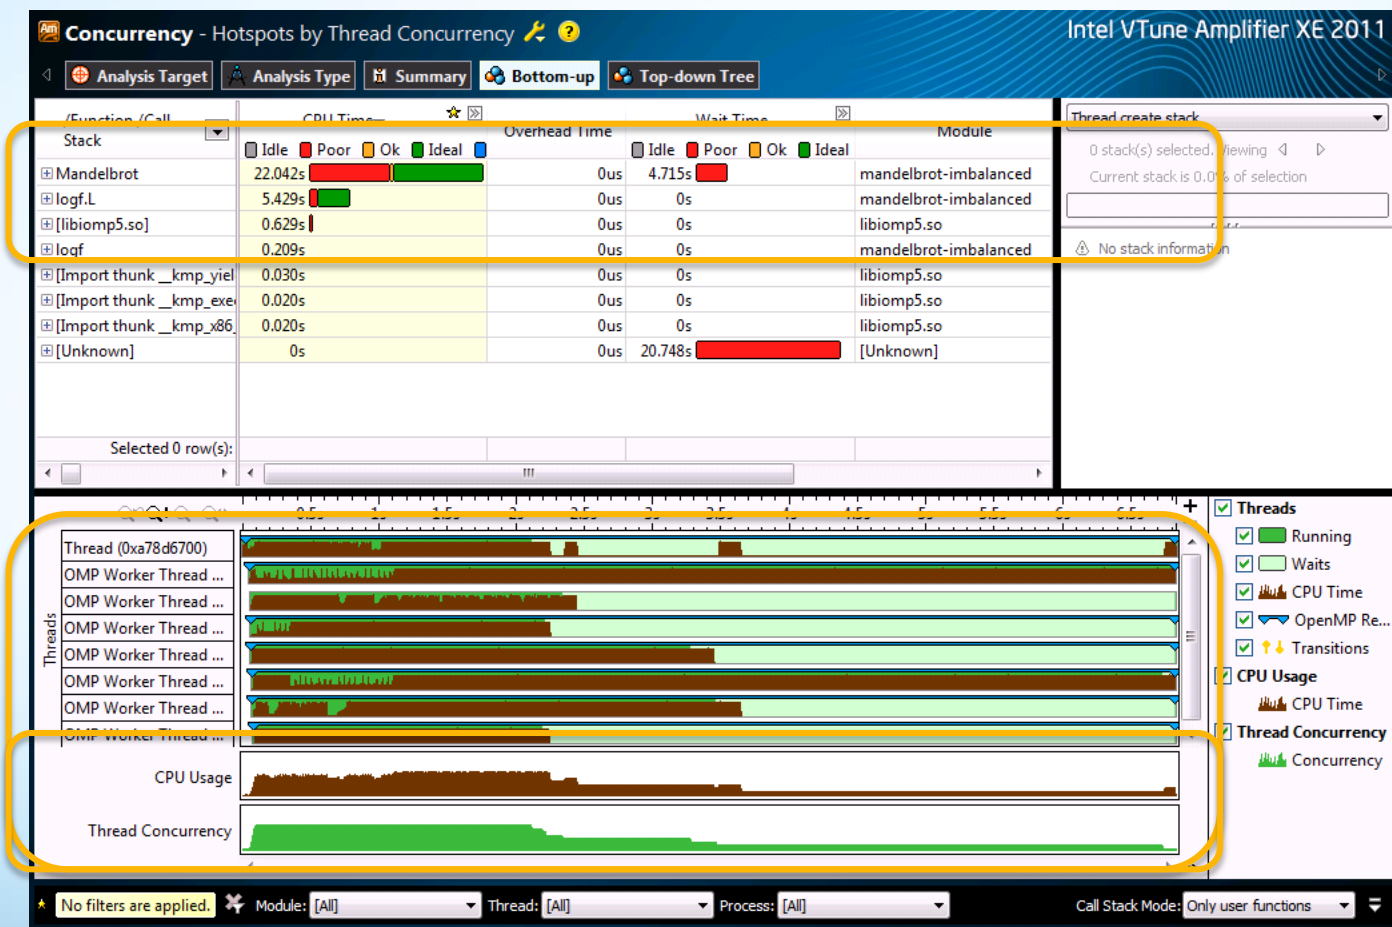
Analysis Types

Concurrency




Hotspots Analysis vs. Concurrency Analysis

- Hotspot Analysis and Concurrency Analysis are similar:



Thread Concurrency Histogram



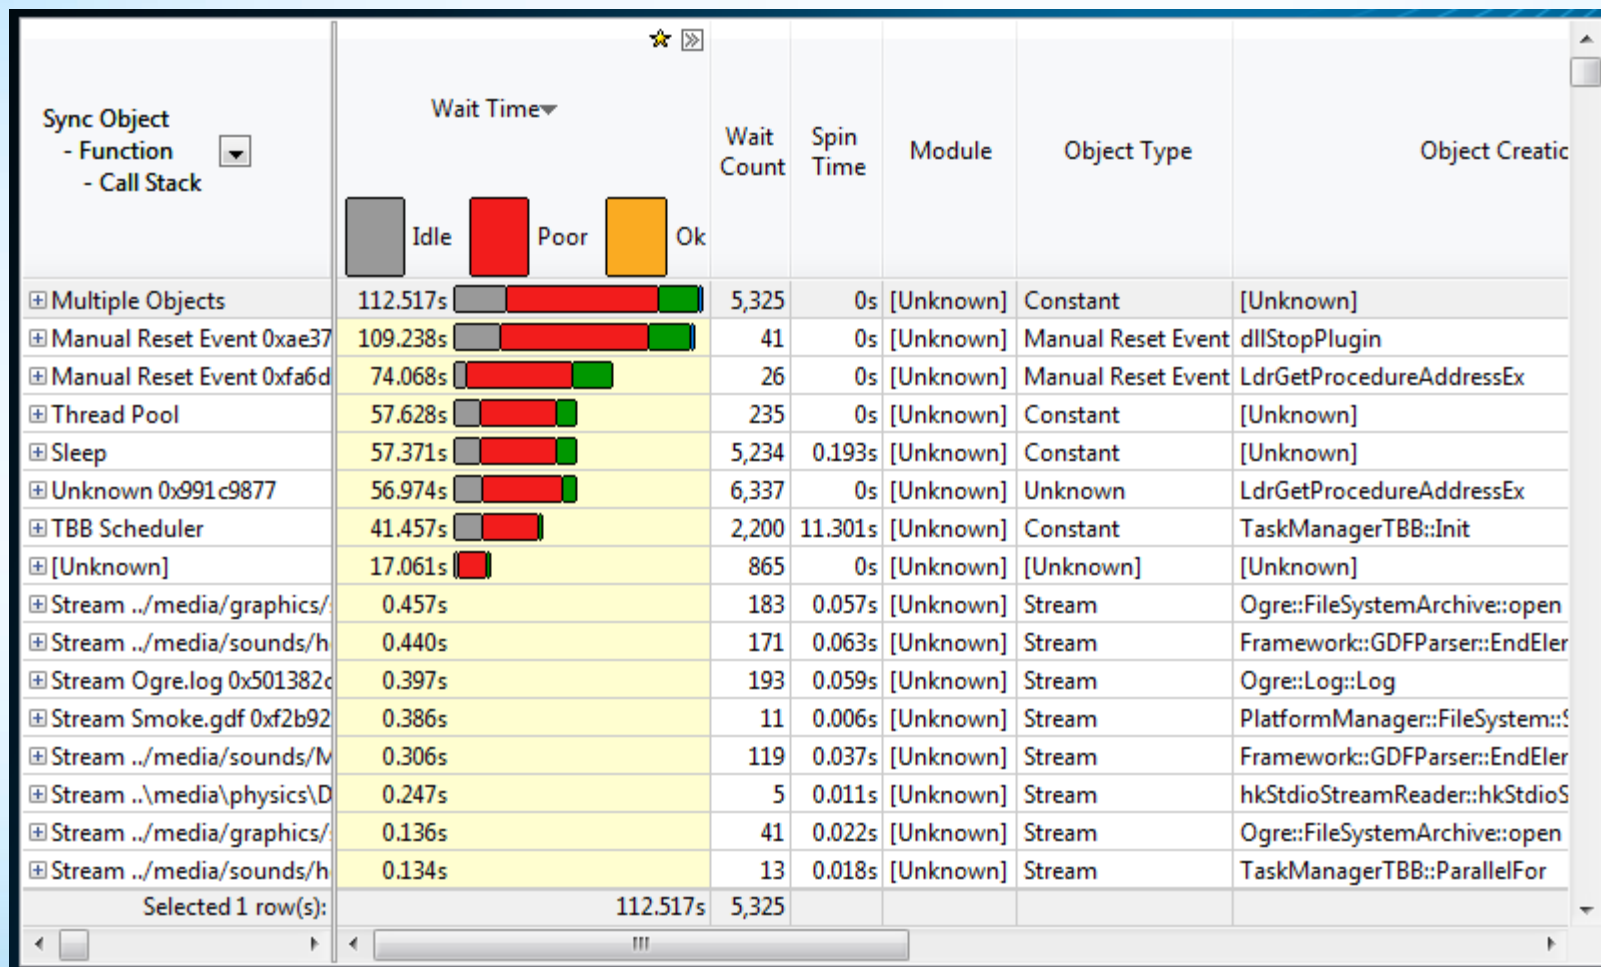
Average Target

Issues in Parallel Applications

- Load imbalance
 - Work distribution is not optimal
 - Some threads are heavily loaded, while others idle
 - Slowest thread determines total speed-up
- Locking issues
 - Locks prohibit threads to concurrently enter code regions
 - Effectively serialize execution
- Parallelization overhead
 - With large no. of threads, data partition get smaller
 - Overhead might get significant (e.g. OpenMP startup time)

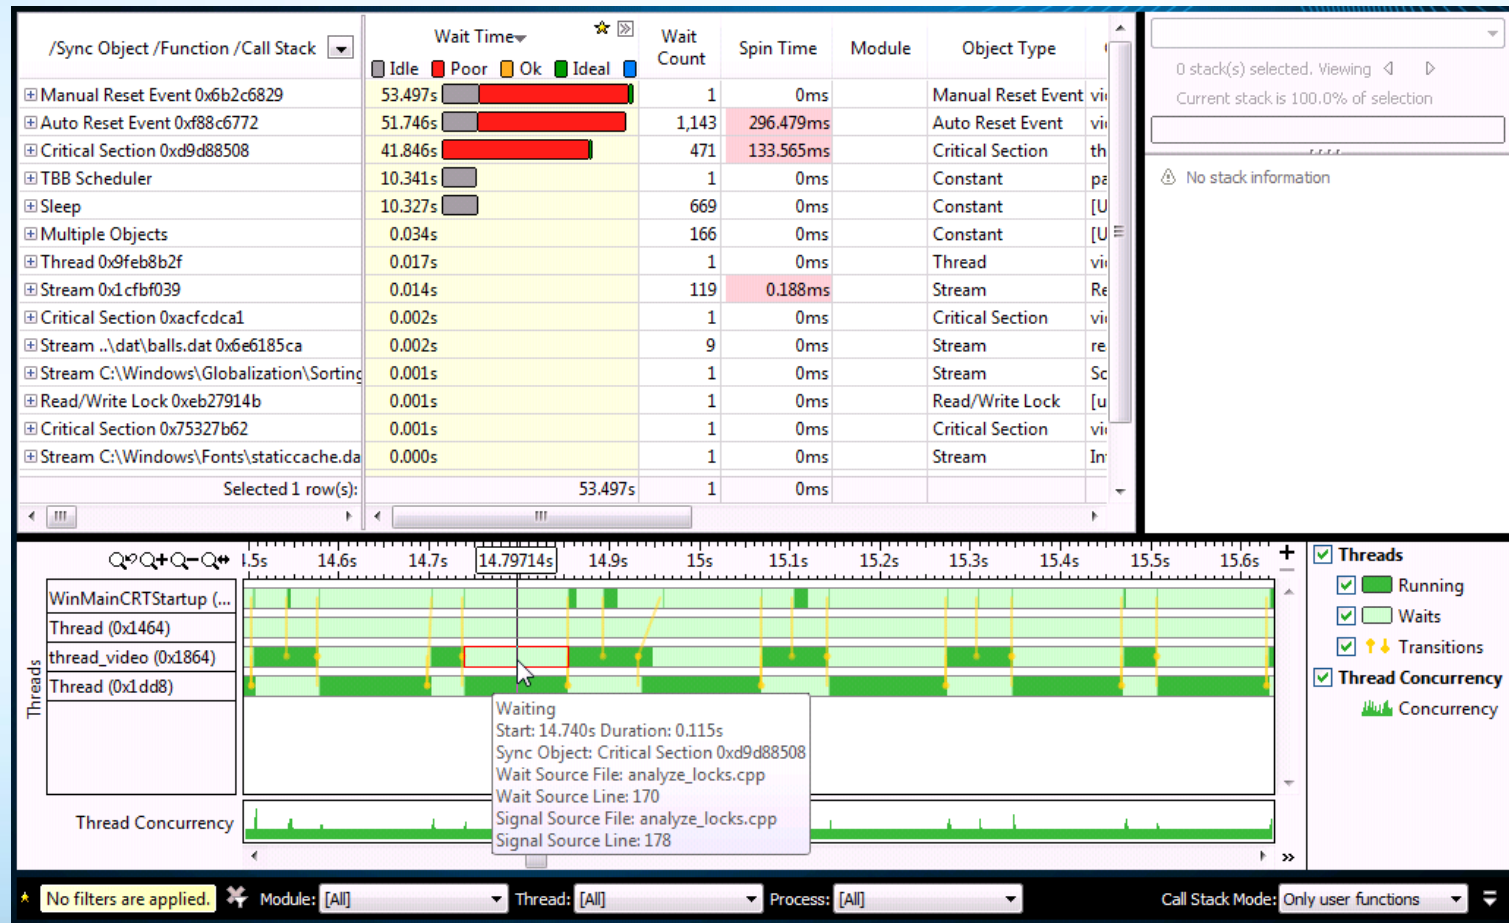
Analysis Types

Lock and Waits



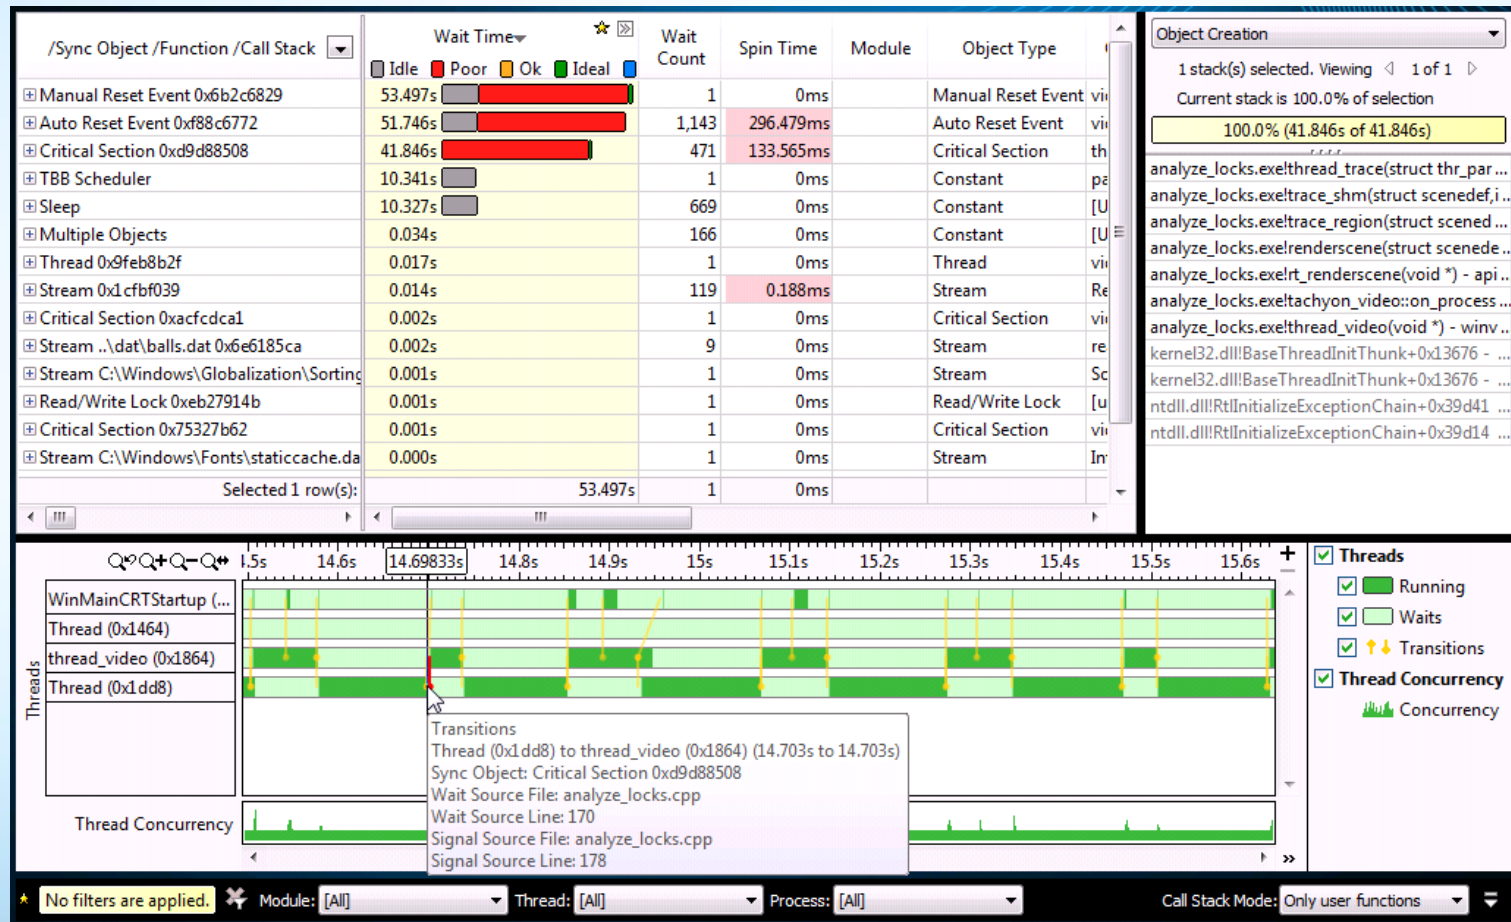
Timeline Visualizes Thread Behavior

- Retrieve additional information about waiting threads



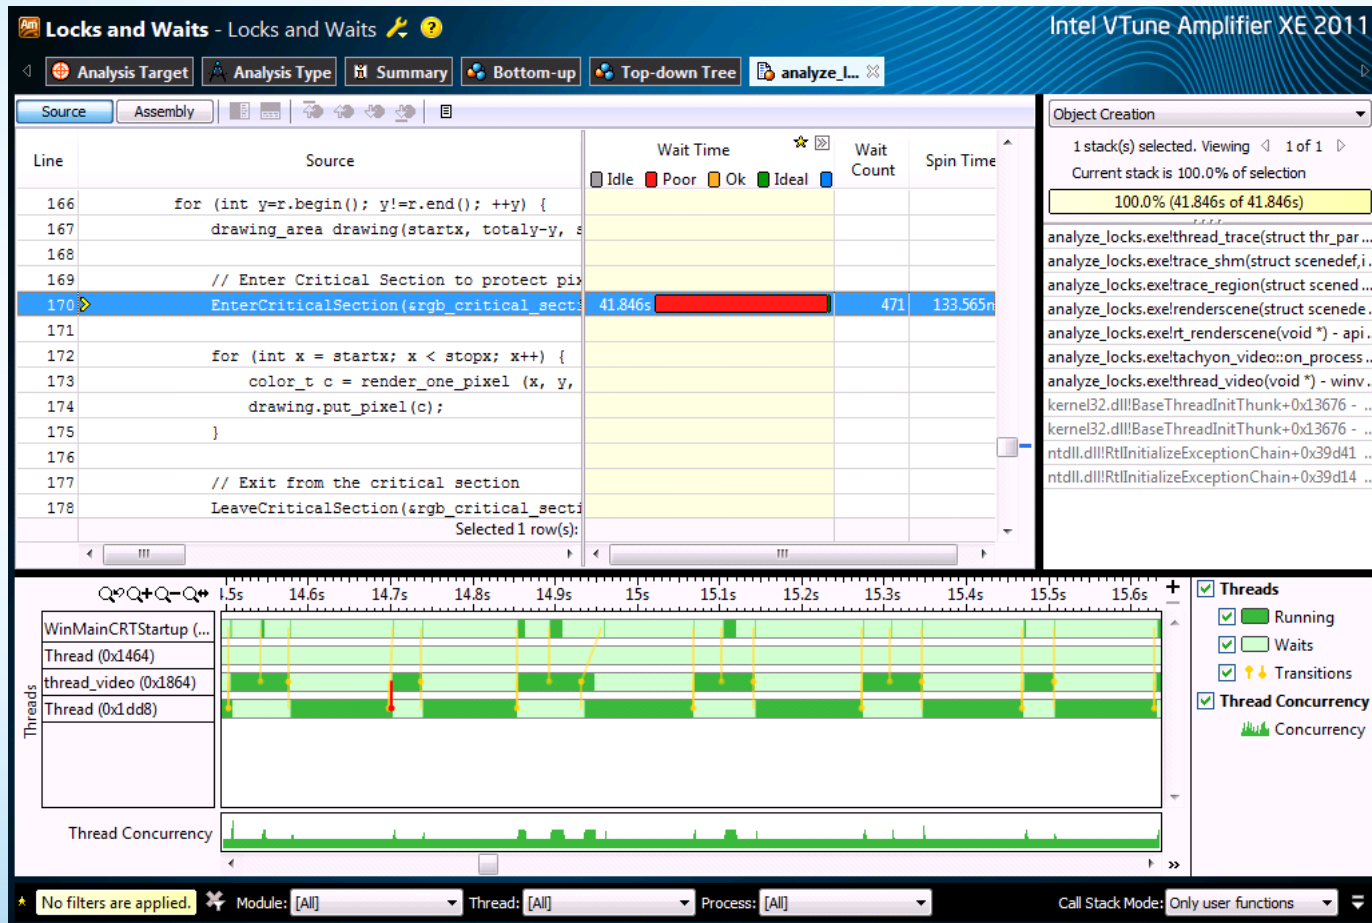
Timeline Visualizes Thread Behavior

- Retrieve additional information on thread transitions



Drilling down into Thread Behavior

- Reveal source code that causes thread transitions



Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Performance Monitoring Unit (PMU)

- **Per-core PMU:**

- Each core provides 2 programmable counters and 1 fixed counter.
- The programmable per-core counters can be configured to investigate front-end/micro-op flow issues, stalls inside a processor core.

- **Uncore PMU:**

- Uncore of the coprocessor has four counters to monitor uncore events
- Can be used to investigate memory behavior and global on-chip issues

Event Based Performance Analysis

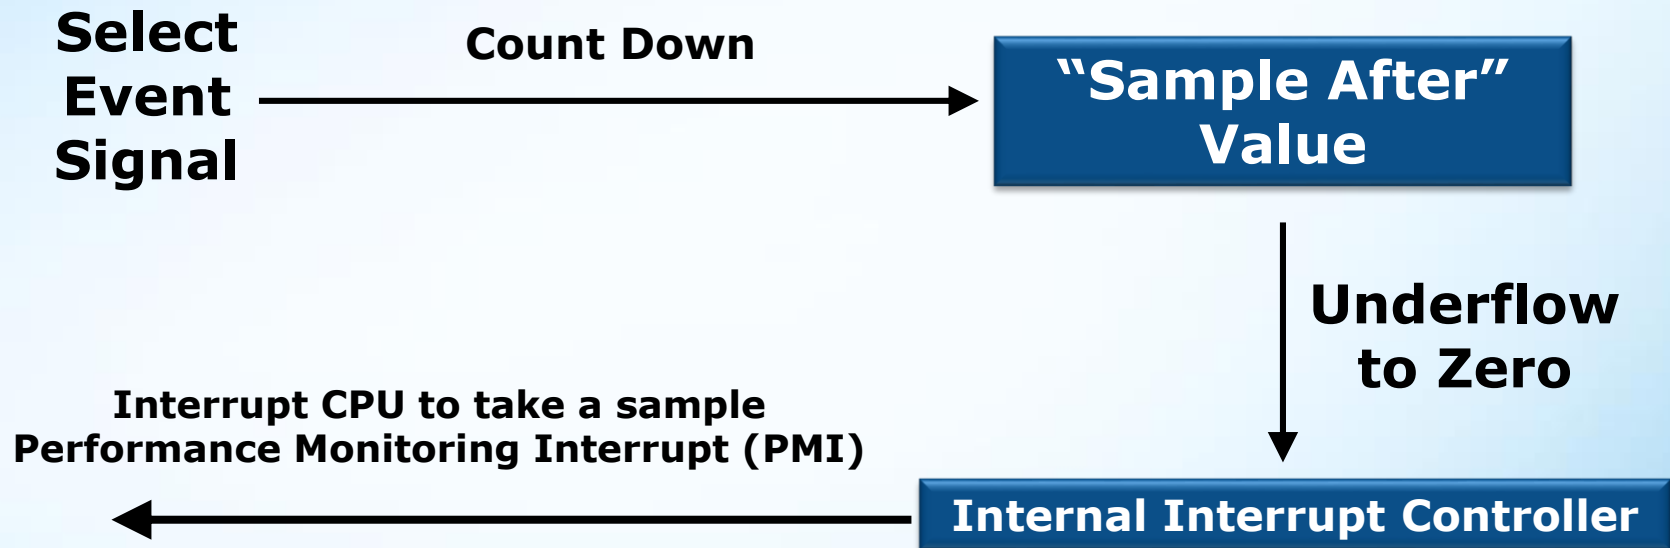
Event Based Sampling(EBS)

- Both architectural and non-architectural processor events can be monitored using sampling and counting technologies

Sampling: Allows to profile all active software on the system, including operating system, device driver, and application software.

- Event-based samples are collected periodically after a specific number of processor events have occurred while the program is running
- The program is interrupted, allowing the interrupt handling driver to collect the Instruction Pointer (IP), load module, thread and process ID's
- Instruction pointer is then used to derive the function and source line number from the debug information created at compile time

How Event Based Sampling (EBS) Works



- A performance counter increments on the CPU every time an event occurs
- A sample of the execution context is recorded every time a performance counter overflows

$$\text{Events} = \text{samples} * \text{sample after value}$$

Native Launch configuration

Target Search Directories

Target type Launch Application

Launch Application
Specify and configure the application executable (target) to analyze. Press F1 for more details.

Application: ssh Browse...

Application parameters: mic0 /tmp/diffusionMIC tiled Modify...

Working directory: /home/michome/rreed/projects/MICtest/perfMIC Browse...

- Application settings:
 - Application: ssh
 - Parameters: mic0 "<app startup>"
 - Working directory: Usually does not matter
 - Don't forget to set search directories under "All files"

Target Search Directories

Search directories for: All files

Search Directories	Search Subdirectories
/home/michome/rreed/projects/MICtest	<input type="checkbox"/>
/lib/firmware/mic	<input checked="" type="checkbox"/>

Application Configuration

File View Help

r0001h r0011h r0021h r0031h New Amplifier XE Result

Choose Analysis Type

Analysis Type

- Access Contention
- Branch Analysis
- Client Analysis
- Core Port Saturation
- Cycles and uOps
- Loop Analysis
- Memory Access
- Port Saturation
- Intel Atom Processor Analysis
 - General Exploration
- Knights Corner Platform Analysis
 - Lightweight Hotspots**
- Power Analysis
 - CPU Sleep States
 - CPU Frequency
- Custom Analysis
 - Branch Mispredicts
 - Cache Misses
 - Execution Stalls
 - False Sharing

Knights Corner Platform - Lightweight Hotspots

Identify your most time-consuming source code. Unlike Hotspots, Lightweight Hotspots has lower overhead because it does not collect stack information. It can also be used to sample all processes on a system. This analysis type uses hardware event-based sampling collection. Press F1 for more details.

List of MIC cards (e.g. 0,1,2,3):

Details

Events configured for CPU: Intel(R) Xeon(R) / Core i7 980X Processor

NOTE: For analysis purposes, Intel VTune Amplifier XE 2013 may adjust the Sample After values in the table below by a multiplier. The multiplier depends on the value of the Duration time estimate option specified in the Project Properties dialog.

Event Name	Sample After	Event De
CPU_CLK_UNHALTED	10000000	Number of cycles during which
INSTRUCTIONS_EXECUTED	10000000	Number of instructions execut

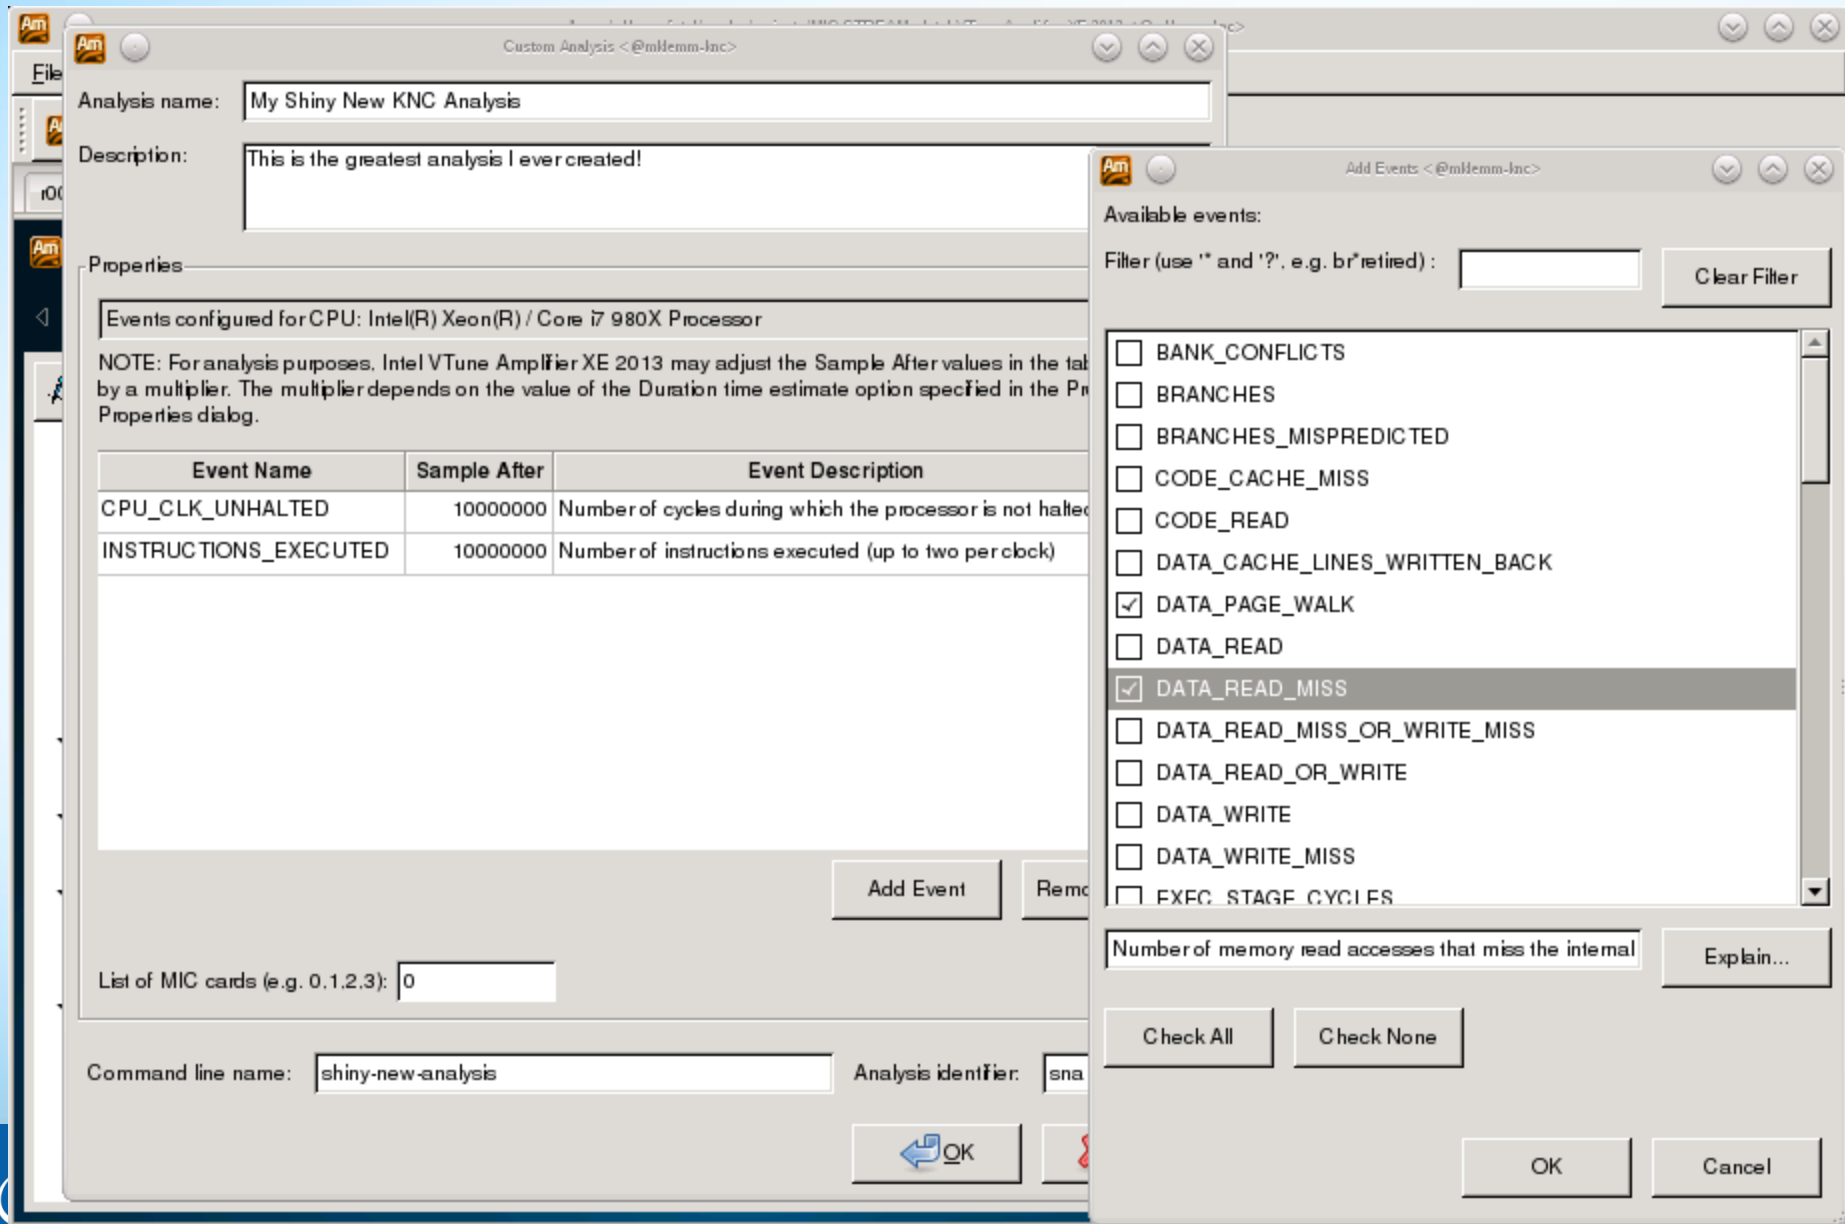
Start

Start Paused

Project Properties

Command Line...

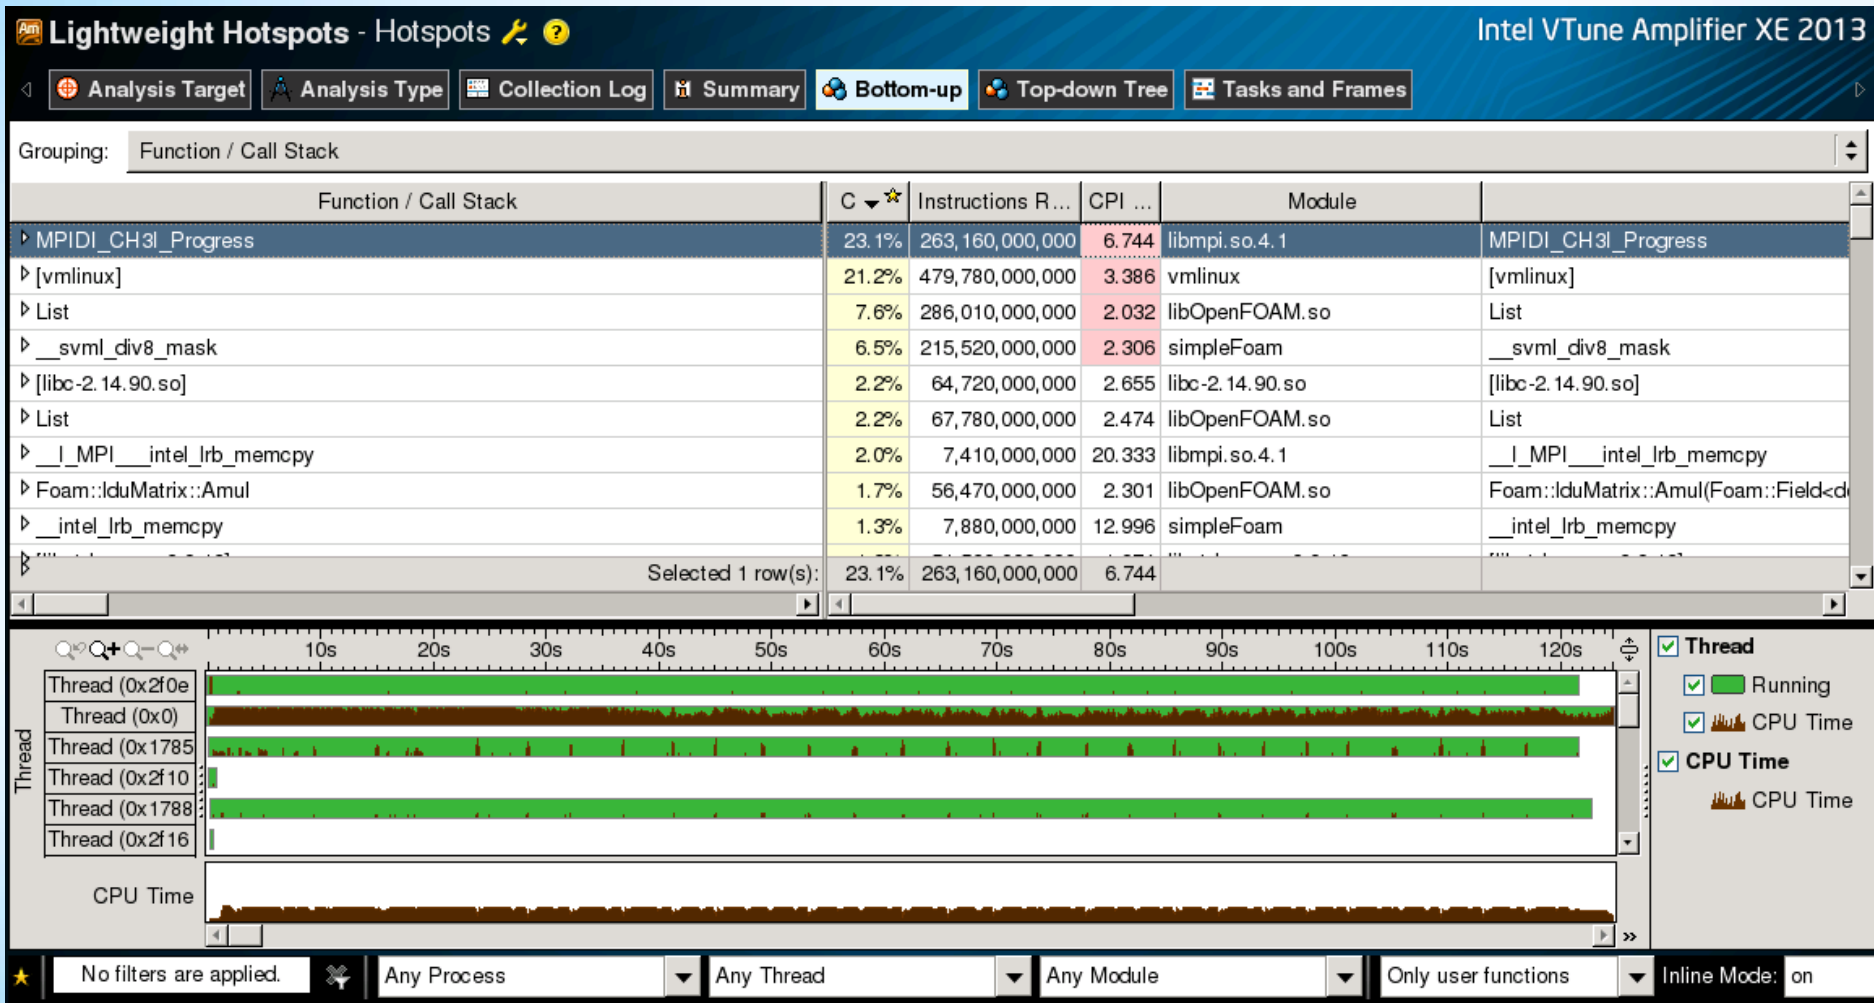
Configuring a User-defined Analysis



Some useful events and metrics

Scenario	Event name(s)
Wall-clock profiling	CPU_CLK_UNHALTED, INSTRUCTIONS_EXECUTED (or EXEC_STAGE_CYCLES)
Main memory bandwidth	L2_DATA_READ_MISS_MEM_FILL, L2_DATA_WRITE_MISS_MEM_FILL
L1 Cache misses	DATA_READ_MISS_OR_WRITE_MISS
TLB misses and page faults	DATA_PAGE_WALK, LONG_DATA_PAGE_WALK, DATA_PAGE_FAULT
Vectorized code execution	VPU_INSTRUCTIONS_EXECUTED, VPU_ELEMENTS_ACTIVE
Various hazards	BRANCHES_MISPREDICTED
Cycles per instruction	CPU_CLK_UNHALTED / INSTRUCTIONS_EXECUTED
Memory Bandwidth (used by all cores at once)	$(L2_DATA_READ_MISS_MEM_FILL + L2_DATA_WRITE_MISS_MEM_FILL) * 64 /$ $CPU_CLK_UNHALTED / \text{Frequency}$

Example: Hotspots of OpenFOAM



Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Data Layout – Common Layouts

Array-of-Structs (AoS)

x	y	z	x	y	z
x	y	z	x	y	z
x	y	z	x	y	z

- **Pros:**
Good locality of {x, y, z}.
1 memory stream.
- **Cons:**
Potential for gather/scatter.

Struct-of-Arrays (SoA)

x	x	x	x	x	x
y	y	y	y	y	y
z	z	z	z	z	z

- **Pros:**
Contiguous load/store.
- **Cons:**
Poor locality of {x, y, z}.
3 memory streams.

Hybrid (AoSoA)

x	x	y	y	z	z
x	x	y	y	z	z
x	x	y	y	z	z

- **Pros:**
Contiguous load/store.
1 memory stream.
- **Cons:**
Not a “normal” layout.

Data Layout – Why It's Important


- Instruction-Level

- Hardware is optimized for contiguous loads/stores.
- Support for non-contiguous accesses differs with hardware.
(e.g. AVX2/KNC gather)

- Memory-Level

- Contiguous memory accesses are cache-friendly.
- Number of memory streams can place pressure on prefetchers.

Data Alignment – Why It's Important



Aligned Load

- Address is aligned.
- One cache line.
- One instruction.

Unaligned Load

- Address is not aligned.
- Potentially multiple cache lines.
- Potentially multiple instructions.

Data Alignment – Why It's Important

- Cache Associativity

- L1 and L2 on Knights Corner are 8-way associative.
 - L1: 8 cache lines that are $32\text{KB}/8 = 4\text{KB}$ apart.
 - L2: 8 cache lines that are $512\text{KB}/8 = 64\text{KB}$ apart.

- Set Conflicts

- Occur when references are a multiple of 4K (L1) or 64K (L2) apart.
- Look for high cache miss rate, even though working set < cache capacity.
- Solution is to pad arrays appropriately.

Data Alignment – Sample Applications

1) Align memory

- `_mm_malloc(bytes, 64)` / `!dir$ attributes align:64`

2) Access Memory in an aligned way

- `for (i = 0; i < N; i++) { array[i] ... }`

3) Tell the compiler!

- `#pragma vector aligned` / `!dir$ vector aligned`
- `__assume_aligned(p, 16)` / `!dir$ assume_aligned (p, 16)`
- `__assume(i % 16 == 0)` / `!dir$ assume (mod(i, 16) .eq. 0)`

Data Alignment – Real-life Applications

0	1	2	3	4	5	6	7
8	9	...					


 Data

Data Alignment – Real-life Applications


0	1	2	3	4	5	6	7
8	9	...					

 Data

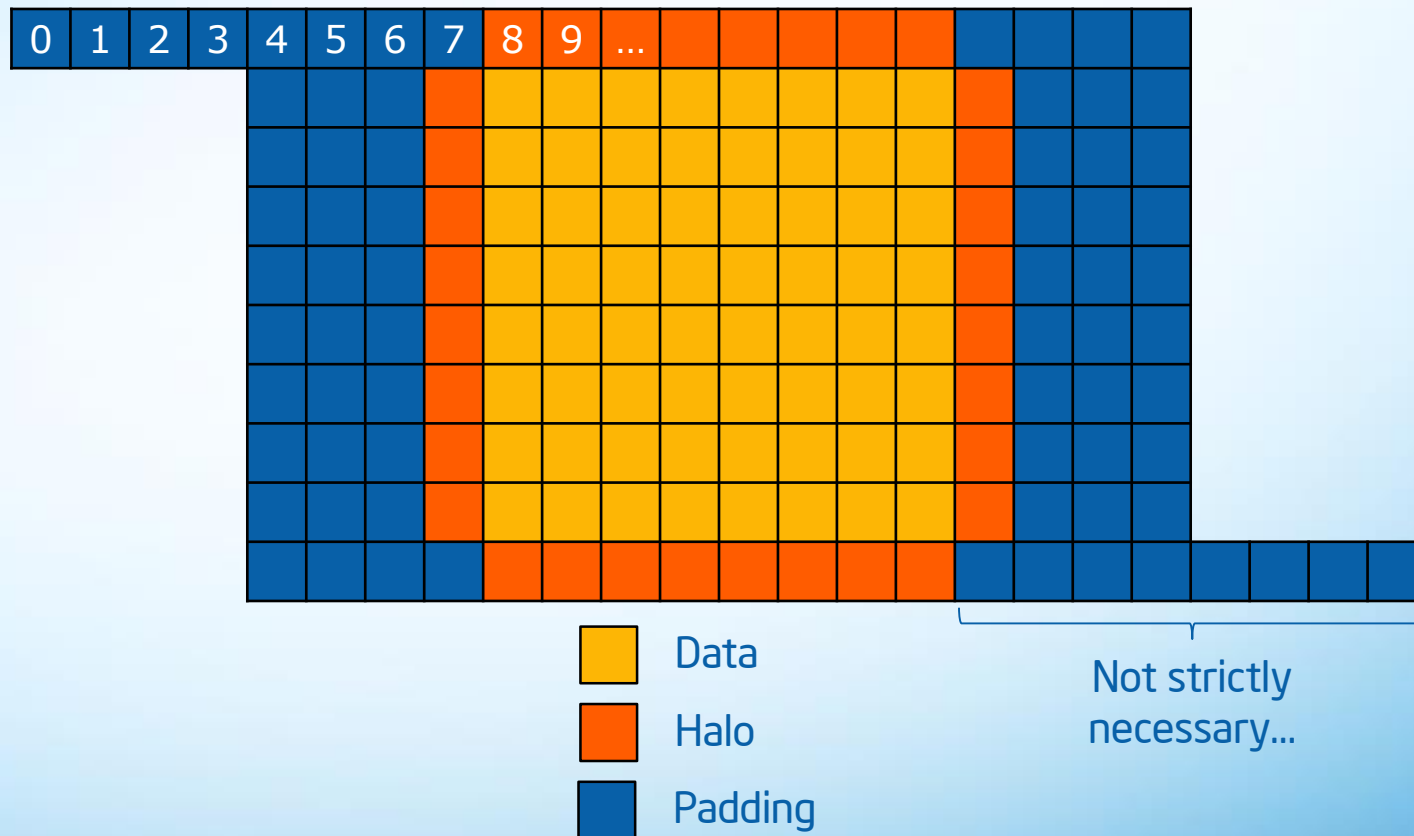
Data Alignment – Real-life Applications




Data Alignment – Real-life Applications



Data Alignment – Real-life Applications



Data Alignment – Real-life Applications



Example: saxpy

```
void saxpy(float *y, float *x, float a, int n) {  
    int i;  
    for (i = 0; i < n; ++i){  
        y[i] = a*x[i] + y[i];  
    }  
}
```

The Effect of Data Alignment on saxpy

Unaligned

```
vloadunpackld (%rsi,%r14,4),%zmm1
vprefetch1 0x400(%rsi,%r14,4)
vloadunpackld 0x40(%rsi,%r14,4),%zmm2
vprefetch0 0x200(%rsi,%r14,4)
vloadunpackhd 0x40(%rsi,%r14,4),%zmm1
vprefetche1 0x400(%rdi,%r14,4)
vloadunpackhd 0x80(%rsi,%r14,4),%zmm2
vprefetch0 0x200(%rdi,%r14,4)
vfmadd213ps (%rdi,%r14,4),%zmm0,%zmm1
vprefetch1 0x440(%rsi,%r14,4)
vfmadd213ps 0x40(%rdi,%r14,4),%zmm0,%zmm2
vprefetch0 0x240(%rsi,%r14,4)
vmovaps %zmm1, (%rdi,%r14,4)
vprefetche1 0x440(%rdi,%r14,4)
vmovaps %zmm2, 0x40(%rdi,%r14,4)
vprefetch0 0x240(%rdi,%r14,4)
```

Aligned

```
vmovaps (%rsi,%r9,4),%zmm2
vprefetch1 0x400(%rsi,%r9,4)
vmovaps 0x40(%rsi,%r9,4),%zmm3
vprefetch0 0x200(%rsi,%r9,4)
vfmadd213ps (%rdi,%r9,4),%zmm0,%zmm2
vprefetche1 0x400(%rdi,%r9,4)
vfmadd213ps 0x40(%rdi,%r9,4),%zmm0,%zmm3
vprefetch0 0x200(%rdi,%r9,4)
vmovaps %zmm2, (%rdi,%r9,4)
vprefetch1 0x440(%rsi,%r9,4)
vmovaps %zmm3, 0x40(%rdi,%r9,4)
vprefetch0 0x240(%rsi,%r9,4)
vprefetche1 0x440(%rdi,%r9,4)
mov %al,%al
vprefetch0 0x240(%rdi,%r9,4)
```

Data Alignment – Host

- AVX instructions are less sensitive to unaligned accesses.
- Still makes sense to align data on the host
 - Cache aligned
 - 4K page aligned
- Data alignment ensures that DMA transfers work best when offloading data to the target
 - Ideal alignment: 4K pages

Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Loop Optimizations

- SIMD works best, if
 - data is properly aligned
 - data is stored in large portions of stride 1
 - loops have a high trip count
- Vectorization potential
 - Average utilization of SIMD register for a loop
 - $VP = (N/vp) / \text{ceil}(N/vl)$
- Vectorization potential can be measured
 - % vector instructions:
`VPU_INSTRUCTIONS_EXECUTED/INSTRUCTIONS_EXECUTED`
 - Avg. % elements used per vector:
`VPU_ELEMENTS_ACTIVE/INSTRUCTIONS_EXECUTED`

Example Loop: NWChem

```
do p4=1,p4d
do p5=1,p5d
do p6=1,p6d
do h1=1,h1d
do h7=1,h7d
do h2=1,h2d
do h3=1,h3d
triplex(h3,h2,h1,p6,p5,p4) =
    triplex(h3,h2,h1,p6,p5,p4) -
        t2sub(h7,p4,p5,h1)*
        v2sub(h3,h2,p6,h7)
end do
end do
end do
end do
end do
end do
```

- Each loop runs on the order of 10 to 25 iterations
- Example: trip count 19
- Vectorization potential
 - SIMD length: 8
 - 2 registers to cover 16 iterations
 - remainder loop: 3 trips

80%

Example Loop: NWChem - Optimized

```
do p4=1,p4d
do p5=1,p5d
do p6=1,p6d
do h1=1,h1d
do h7=1,h7d
do h3h2=1,h2d*h3d
triplex(h3h2,h1,p6,p5,p4) =
  triplex(h3h2,h1,p6,p5,p4) -
    t2sub(h7,p4,p5,h1)*
    v2sub(h3h2,p6,h7)
end do
end do
end do
end do
end do
```

- Fused loop has about 1 or 2 orders of magnitudes more trips
- Example: trip count 19
- Vectorization potential
 - SIMD length: 8
 - 45 registers to cover 360 iterations
 - remainder loop: 1 trips

98%

Agenda

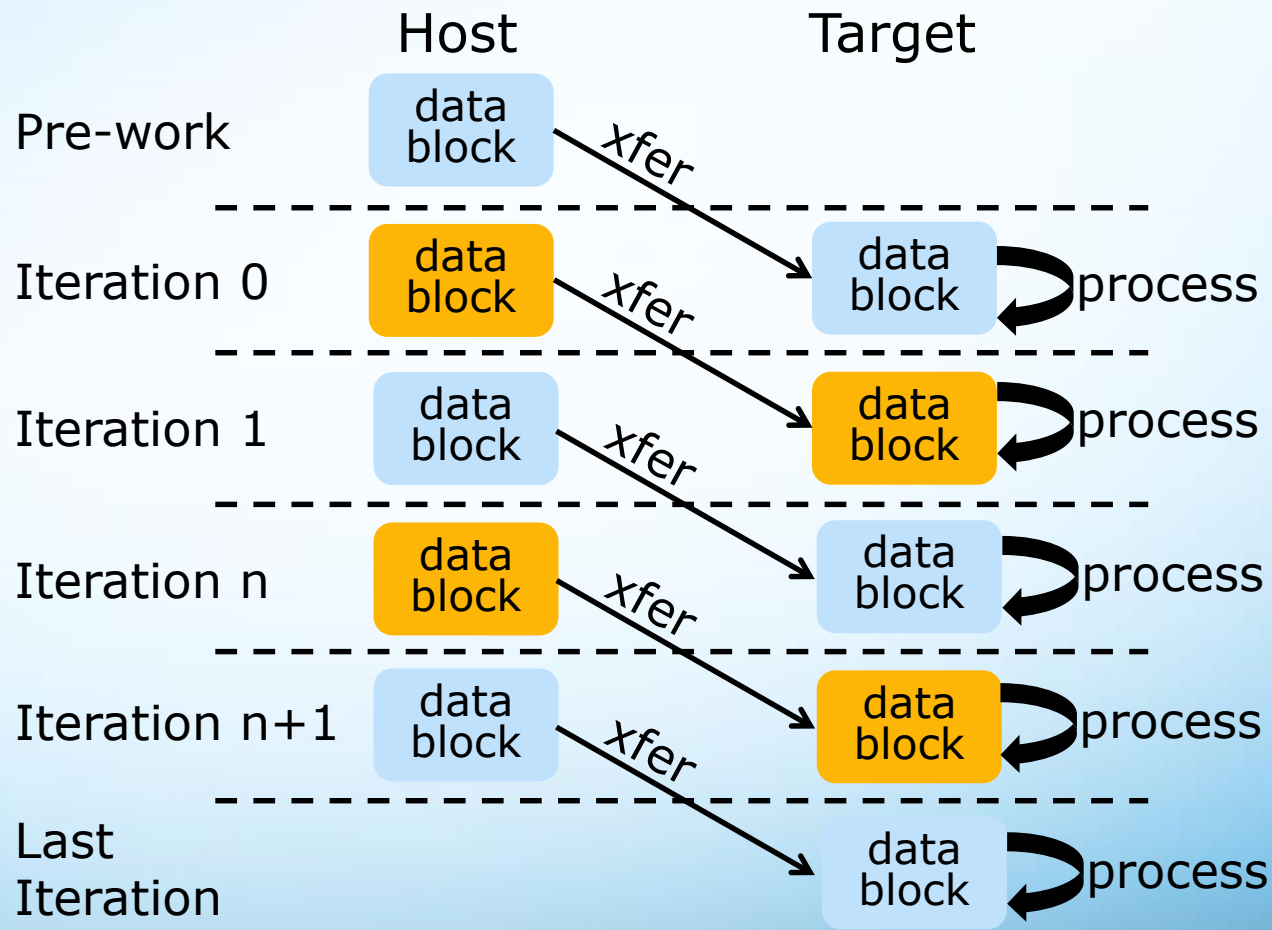
- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Optimizing Data Transfers

- PCIe is a bottleneck compare to memory
 - High latency (on the order of microseconds)
 - Low bandwidth (1 to 2 orders of magnitude less than GDDR)
- Avoid data transfers as much as possible
- Overlap data transfers and computation where possible
- Stay on the target device as long as possible

Double Buffering Example

- Overlap computation and communication
- Generalizes to data domain decomposition



Double Buffering Example - Main

```
int main(int argc, char* argv[]) {  
    int i;  
    double st_time, end_time;  
    double sync_tm, async_in_tm;  
    // Allocate & initialize in1, res1,  
    // in2, res2 on the host
```

Allocate arrays on target:
`alloc_if(1)`

Retain for duration of sample:
`free_if(0)`

in1 and in2 represent 2 separate buffers.

```
#pragma offload_transfer target(mic:0) in(cnt) \  
    nocopy(in1, res1, in2, res2 : length(cnt) \  
    alloc_if(1) free_if(0) )
```

```
do_async_in();
```

```
// Validate results and print timings
```

Free target allocations:
`free_if(1)`

```
#pragma offload_transfer target(mic:0) \  
    nocopy(in1, res1, in2, res2 : length(cnt) \  
    alloc_if(0) free_if(1) )
```

Double Buffering – do_async_in, evens

```
void do_async_in() {  
    float lsum;  
    int i;  
    lsum = 0.0f;
```

Begin an initial transfer of the first dataset for the target to work on. Transfer begins immediately, is non-blocking, and will signal when complete.

```
#pragma offload_transfer target(mic:0) \  
    in(in1 : length(cnt) alloc_if(0) free_if(0) ) signal(in1)
```

```
for (i=0; i < iter; i++) {  
    if (i%2 == 0) {
```

For even loop iterations (except the final), first start another non-blocking transfer of the next dataset.

```
#pragma offload_transfer target(mic:0) if(i!=iter-1) \  
    in(in2 : length(cnt) alloc_if(0) free_if(0) ) \  
    signal(in2)
```

```
#pragma offload target(mic:0) nocopy(in1) wait(in1) \  
    out(res1 : length(cnt) alloc_if(0) free_if(0) )  
{  
    compute(in1, res1);  
}
```

While that transfer progresses, process the previous dataset through the compute() function, first waiting for its transfer to complete, then return a result. Execution on the host waits for the function to return.

```
    lsum = lsum + sum_array(res1);  
}
```

Double Buffering – do_async_in, odds

```
else {
```

```
#pragma offload_transfer target(mic:0) if(i!=iter-1) \  
    in(in1 : length(cnt) alloc_if(0) free_if(0) ) \  
    signal(in1)
```

For odd iterations (except the final), work on the other buffer. Start another non-blocking transfer.

```
#pragma offload target(mic:0) nocopy(in2) wait(in2) \  
    out(res2 : length(cnt) alloc_if(0) free_if(0) )  
{  
    compute13(in2, res2);  
}
```

Offload the compute function, and host waits for the result. Repeat, alternating between the buffers (in1 & in2).

```
    lsum = lsum + sum_array(res2);  
}  
}  
async_in_sum = lsum / (float) iter;  
}
```


Agenda

- Performance Tuning Methodology
- Intel® VTune™ Amplifier XE
 - Fundamental Analysis: Hotspots
 - Finding Issues in Parallel Applications
 - Using the Performance Monitoring Unit
- Optimization Strategies
 - Data Layout and Alignment
 - Vectorization Potential
 - Optimizing Data Transfers
 - Miscellaneous Tuning Hints

Miscellaneous Tuning Hints

- Check if data transfer is necessary
 - Keep data as long on the device as possible
 - Avoid transferring scratch pad memory
- Use 2MB pages for offload data transfers
 - (May) result in less page faults
 - (May) improve utilization of the TLB cache
 - (May) increase time needed to initialize pages
- Controlled by environment variable:
`export MIC_USE_2MB_BUFFERS=16k`

Miscellaneous Tuning Hints

- Use 2MB pages for dynamic memory
- Reserve 2MB pages in the kernel:
 - Remember this leaves less memory available for other programs!
 - Example (reserves 128 2MB pages):
echo "128" > /proc/sys/vm/nr_hugepages
- Use mmap to request allocation of 2MB pages:

```
#include <sys/mman.h>
```

```
void *p;
```

```
size_t s = 64*2*1024*1024; // 128MB
```

```
...
```

```
p = mmap(0, s, PROT_READ|PROT_WRITE,  
        MAP_ANONYMOUS|MAP_PRIVATE|MAP_HUGETLB, -1, 0);
```

```
if (p == MAP_FAILED)  
    perror("mmap failed");
```

Miscellaneous Tuning Hints

- Always use thread affinity on the coprocessor
 - Avoid expensive thread migration between cores
 - Example:

```
export MIC_ENV_PREFIX=MIC
export MIC_KMP_AFFINITY=compact / scatter / balanced
```
 - Affinity policies:
 - compact: keep threads close
 - scatter: distribute threads as far apart as possible
 - balanced: create *compact* groups and *scatter* these
 - Simpler approach:

```
export KMP_PLACE_THREADS=60c,3t,0o
```

Miscellaneous Tuning Hints

- Do not overload the coprocessor
 - Multiple offloads from different host MPI ranks
 - Multiple offloads from different host threads
- Partition the device:
 - Rank 0: export MIC_PLACE_THREADS=30c,3t,0o
 - Rank 1: export MIC_PLACE_THREADS=30c,3t,30o
- May help to improve relative performance
 - Two offloads will half the compute load each might be more beneficial
 - Less synchronization / threading overhead

Summary

- Use profiling tools to find hotspots and offload candidates
- Use a performance analysis tool to find threading issues
- Optimize offloaded code
 - Data layout and alignment
 - Vectorization potential

